

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
plt.style.use('ggplot')
%matplotlib inline
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, f1_score, classification_
```

In [2]:

```
data_train=pd.read_csv('train.csv')
data_test=pd.read_csv('test.csv')
```

In [3]:

```
print('Shape of train dataset is {}'.format(data_train.shape))
print('Shape of test dataset is {}'.format(data_test.shape))
```

```
Shape of train dataset is (9557, 143)
Shape of test dataset is (23856, 142)
```

## Identify the output variable.

In [4]:

```
target_col = ""
for i in data_train.columns:
    if i not in data_test.columns:
        target_col = i
        print("Our output variable is {}".format(i))
        break
```

Our output variable is Target

## Understand the type of data.

In [5]:

```
data_train.dtypes.value_counts()
```

Out[5]:

```
int64      130
float64      8
object      5
dtype: int64
```

In [6]:

```
data_train.select_dtypes(include='object')
```

Out[6]:

|      | Id           | idhogar   | dependency | edjefe | edjefa |
|------|--------------|-----------|------------|--------|--------|
| 0    | ID_279628684 | 21eb7fcc1 | no         | 10     | no     |
| 1    | ID_f29eb3ddd | 0e5d7a658 | 8          | 12     | no     |
| 2    | ID_68de51c94 | 2c7317ea8 | 8          | no     | 11     |
| 3    | ID_d671db89c | 2b58d945f | yes        | 11     | no     |
| 4    | ID_d56d6f5f5 | 2b58d945f | yes        | 11     | no     |
| ...  | ...          | ...       | ...        | ...    | ...    |
| 9552 | ID_d45ae367d | d6c086aa3 | .25        | 9      | no     |
| 9553 | ID_c94744e07 | d6c086aa3 | .25        | 9      | no     |
| 9554 | ID_85fc658f8 | d6c086aa3 | .25        | 9      | no     |
| 9555 | ID_ced540c61 | d6c086aa3 | .25        | 9      | no     |
| 9556 | ID_a38c64491 | d6c086aa3 | .25        | 9      | no     |

9557 rows × 5 columns

## Check if there are any biases in your dataset.

In [7]:

```
(data_train[target_col].value_counts() / data_train.shape[0]) * 100
```

Out[7]:

```
4    62.739353
2    16.710265
3    12.650413
1     7.899969
Name: Target, dtype: float64
```

Above data shows that about 62% of people in training dataset fall under category 4 level of poverty

In [8]:

```
fig, axes = plt.subplots(3,2,figsize=(10,15))

sns.histplot(ax=axes[0,0], x=data_train['r4t3'],
             bins= 10, kde = True, cbar = True,
             color = '#A163CF').set(title = "Distribution of 'r4t3'")

sns.histplot(ax=axes[0,1], x=data_train['hogar_total'],
             bins = 10, kde = True, cbar = True,
             color = '#A163CF').set(title = "Distribution of 'hogar_total'")

sns.histplot(ax=axes[1,0], x=data_train['tipovivi3'],
             bins = 10, kde = True, cbar = True,
             color = '#A163CF').set(title= "Distribution of 'tipovivi3'")


sns.histplot(ax=axes[1,1], x=data_train['v2a1'],
             bins=10, kde=True,cbar = True,
             color = '#A163CF').set(title="Distribution of 'v2a1'")

sns.histplot(ax=axes[2,0],x=data_train['v18q'],
             bins=10, kde = True, cbar = True,
             color = '#A163CF').set(title="Distribution of 'v18q'")

sns.histplot(ax=axes[2,1],x=data_train['v18q1'],
             bins=10, kde = True, cbar = True,
             color = '#A163CF').set(title="Distribution of 'v18q1'")
```

Out[8]:

```
[Text(0.5, 1.0, "Distribution of 'v18q1'")]
```



Therefore, variables ('r4t3', 'hogar\_total') have relationship between them. For good result we can use any one of them. Therefore, variables ('tipovivi3', 'v2a1') have relationship between them. For good result we can use any one of them. Therefore, variables ('v18q', 'v18q1') have relationship between them. For good result we can use any one of them. So, *Therefore, there is bias in our dataset.*

**Check whether all members of the house have the same poverty level.**

In [9]:

```
# Group By data wrt to family ID
family_group = data_train.groupby('idhogar')
```

In [10]:

```
# Get the count unique count of Target variable
unique_poverty_status = family_group[target_col].nunique()
# unique_poverty_status[unique_poverty_status > 1].index
any(unique_poverty_status > 1)
```

Out[10]:

True

In [11]:

```
sum(unique_poverty_status > 1), unique_poverty_status.shape[0]
```

Out[11]:

(85, 2988)

Above data shows that there are some families where members are of different poverty category within same family. Out of 2988 Families in dataset about 85 families have shown this behaviour

## Check if there is a house without a family head.

In [12]:

```
num_head_in_family = family_group.parentesco1.sum()
any(num_head_in_family == 0)
```

Out[12]:

True

In [13]:

```
num_head_in_family[num_head_in_family == 0].shape[0], num_head_in_family.shape[0]
```

Out[13]:

(15, 2988)

Above data shows that out 2988 families about 15 families don't have family head in training dataset

## Set the poverty level of the members and the head of the house same in a family.

In [14]:

```
def set_poverty_level_as_head(group):
    head_poverty_level = group[group.parentesco1 == 1].Target
    if not head_poverty_level.empty:
        group.Target = head_poverty_level.values[0]
    return group
```

In [15]:

```
# update training data
data_train = family_group.apply(set_poverty_level_as_head)
```

## Count how many null values are existing in columns.

In [16]:

```
null_col_count = data_train.isnull().sum()
null_col_count[null_col_count != 0]
```

Out[16]:

```
v2a1          6860
v18q1         7342
rez_esc       7928
meaneduc         5
SQBmeaned         5
dtype: int64
```

### Checking null of v2a1

In [17]:

```
data = data_train[data_train['v2a1'].isnull()].head()

columns=['tipovivi1', 'tipovivi2', 'tipovivi3', 'tipovivi4', 'tipovivi5']
data[columns]
```

Out[17]:

|    | tipovivi1 | tipovivi2 | tipovivi3 | tipovivi4 | tipovivi5 |
|----|-----------|-----------|-----------|-----------|-----------|
| 2  | 1         | 0         | 0         | 0         | 0         |
| 13 | 1         | 0         | 0         | 0         | 0         |
| 14 | 1         | 0         | 0         | 0         | 0         |
| 26 | 1         | 0         | 0         | 0         | 0         |
| 32 | 1         | 0         | 0         | 0         | 0         |


In [18]:

```
# Variables indicating home ownership
own_variables = [x for x in data_train if x.startswith('tipo')]

# Plot of the home ownership variables for home missing rent payments
data_train.loc[data_train['v2a1'].isnull(), own_variables].sum().plot.bar(figsize =
                                                                    color = 'orange',
                                                                    edgecolor = 'k', linewidth = 2)

plt.xticks([0, 1, 2, 3, 4],
            ['Owns and Paid Off', 'Owns and Paying', 'Rented', 'Precarious', 'Other'],
            rotation = 20)
plt.title('Home Ownership Status for Households Missing Rent Payments', size = 18);
```

Home Ownership Status for Households Missing Rent Payments



In [19]:

```
data_train['v2a1'].fillna(0,inplace=True)
```

Looking at the above data it makes sense that when the house is fully paid, there will be no monthly rent payment. So, added 0 for all the null values.

### Checking null values of v18q1

In [20]:

```
data_train['v18q1'].value_counts()
```

Out[20]:

```
1.0    1586
2.0     444
3.0     129
4.0      37
5.0      13
6.0       6
Name: v18q1, dtype: int64
```

In [21]:

```
data_train['v18q1'].isna().sum()
```

Out[21]:

```
7342
```

In [22]:

```
data_train['v18q'].value_counts()
```

Out[22]:

```
0    7342
1    2215
Name: v18q, dtype: int64
```

In [23]:

```
data_train['v18q1'].fillna(0,inplace=True)
```

Looking at the above data it makes sense that when owns a tablet column is 0, there will be no number of tablets household owns. So, added 0 for all the null values.

### ***Checking for null values rez\_esc***

Lets look at rez\_esc : Years behind in school

In [24]:

```
data_train[data_train['rez_esc'].notnull()]['age'].describe()
```

Out[24]:

```
count    1629.000000
mean      12.258441
std        3.218325
min        7.000000
25%        9.000000
50%       12.000000
75%       15.000000
max       17.000000
Name: age, dtype: float64
```



From the above , we see that when min age is 7 and max age is 17 for Years, then the 'behind in school' column has a value.

In [25]:

```
data_train.loc[data_train['rez_esc'].isnull()][ 'age' ].describe()
```

Out[25]:

```
count      7928.000000
mean        38.833249
std         20.989486
min          0.000000
25%         24.000000
50%         38.000000
75%         54.000000
max         97.000000
Name: age, dtype: float64
```

In [26]:

```
data_train.loc[(data_train['rez_esc'].isnull() & ((data_train['age'] > 7) & (data_tr
```

Out[26]:

```
count      1.0
mean       10.0
std        NaN
min        10.0
25%        10.0
50%        10.0
75%        10.0
max        10.0
Name: age, dtype: float64
```

There is one value that has Null for the 'behind in school' column with age between 7 and 17

In [27]:

```
data_train[(data_train['age'] ==10) & data_train['rez_esc'].isnull()].head()
data_train[(data_train['Id'] == 'ID_f012e4242')].head()
```

Out[27]:

|      | Id           | v2a1     | hacdor | rooms | hacapo | v14a | refrig | v18q | v18q1 | r4h1 | ... | SQE |
|------|--------------|----------|--------|-------|--------|------|--------|------|-------|------|-----|-----|
| 2514 | ID_f012e4242 | 160000.0 | 0      | 6     | 0      | 1    | 1      | 1    | 1.0   | 0    | ... |     |

1 rows × 143 columns

There is only one member in household for the member with age 10 and who is 'behind in school'. This explains why the member is behind in school.

In [28]:

```
for data in [data_train, data_test]:
    data['rez_esc'].fillna(value=0, inplace=True)
data_train[['rez_esc']].isnull().sum()
```

Out[28]:

```
rez_esc      0
dtype: int64
```

### Checking for null values meaneduc

In [29]:

```
data = data_train[data_train['meaneduc'].isnull()].head()

columns=['edjefe','edjefa','instlevel1','instlevel2']
data[columns][data[columns]['instlevel1']>0].describe()
```

Out[29]:

|       | instlevel1 | instlevel2 |
|-------|------------|------------|
| count | 0.0        | 0.0        |
| mean  | NaN        | NaN        |
| std   | NaN        | NaN        |
| min   | NaN        | NaN        |
| 25%   | NaN        | NaN        |
| 50%   | NaN        | NaN        |
| 75%   | NaN        | NaN        |
| max   | NaN        | NaN        |

In [30]:

```
for data in [data_train, data_test]:
    data['meaneduc'].fillna(value=0, inplace=True)
data_train[['meaneduc']].isnull().sum()
```

Out[30]:

```
meaneduc      0
dtype: int64
```

### Checking for null values SQBmeaned

In [31]:

```
for data in [data_train, data_test]:  
    data['SQBmeaned'].fillna(value=0, inplace=True)  
data_train[['SQBmeaned']].isnull().sum()
```

Out[31]:

```
SQBmeaned      0  
dtype: int64
```

We find that SQBmeaned is null when no level of education is 0

## Remove null value rows of the target variable.

In [32]:

```
data_train[target_col].isnull().sum()
```

Out[32]:

```
0
```

No Null values found in train data set for Target column

## Predict the accuracy using random forest classifier. (Try other models too)

### *Data Processing*

#### Categorise alpha numeric column

### *Dependency*

In [33]:

```
# Dependency: # of members less than 19 or more than 64 / # of members b/w 19 and 64
data_train['dependency'].value_counts()
```

Out[33]:

```
yes          2192
no           1747
.5           1497
2             730
1.5           713
.33333334     598
.66666669     487
8             378
.25           260
3             236
4             100
.75            98
.2             90
.40000001     84
1.33333334     84
2.5            77
5              24
1.25           18
```

Since it is not clear if how much does a yes quantify to proceeding with assumption of 1 for `yes` and 0 for `no`

In [34]:

```
def map(i):
    if i=='yes':
        return(float(1))
    elif i=='no':
        return(float(0))
    else:
        return(float(i))
```

In [35]:

```
data_train['dependency']=data_train['dependency'].apply(map)
```

In [36]:

```
data_train.dependency.dtypes
```

Out[36]:

```
dtype('float64')
```

In [37]:

```
data_train['dependency'].isnull().sum()
```

Out[37]:

```
0
```

**edjefe & edjefa**

edjefe and edjefa features are not clear and wasn't able to find proper definition so dropping them.

In [38]:

```
data_train = data_train.drop(columns = ['edjefe', 'edjefa'])
```

### ***train & test split***

In [39]:

```
X = data_train.drop(columns=[target_col])
```

In [40]:

```
y = data_train[[target_col]]
```

In [41]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    random_state=42,
                                                    train_size=0.8)
```

In [42]:

```
X_train.shape
```

Out[42]:

```
(7645, 140)
```

In [43]:

```
X_test.shape
```

Out[43]:

```
(1912, 140)
```

## **EDA**

### ***Output variable distribution***

In [44]:

```
target_counts=(y_train[target_col].value_counts() / y_train.shape[0]) * 100
target_counts
```

Out[44]:

```
4    63.047744
2    16.429039
3    12.609549
1     7.913669
Name: Target, dtype: float64
```

In [45]:

```
target_counts.plot.bar(figsize = (8, 6),linewidth = 2,edgecolor = 'k',title="Target
```

Out[45]:

<AxesSubplot:title={'center':'Target vs Total\_Count'}>



**Output variable trends with raw data available**

In [46]:

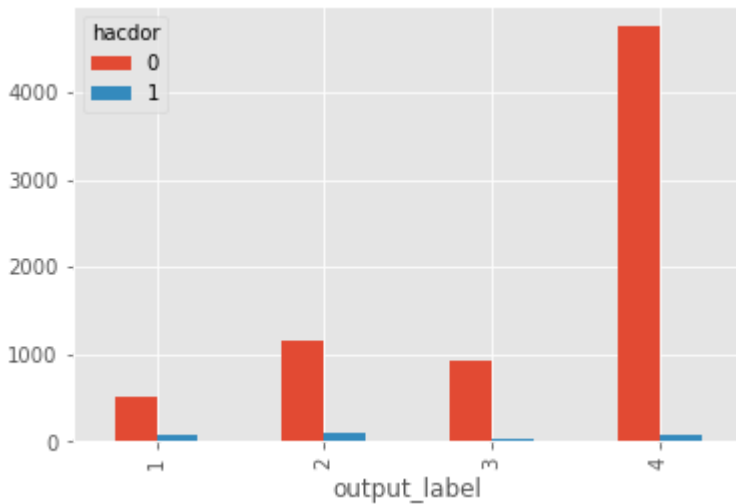
```
x_train['output_label'] = y_train.values
```

In [47]:

```
X_train.groupby(['output_label', 'hacdor']).Id.count().unstack().plot(kind='bar')
```

Out[47]:

&lt;AxesSubplot:xlabel='output\_label'&gt;




Hacdor doesn't seem to provide valuable insight about a person being in particular poverty category or not except for the fact that already biased data showing that if they don't have overcrowded bedroom they are likely to be in category 4, will have to revisit about viability of this feature later

In [48]:

```
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(10, 4))
room_distribution = X_train['rooms'].value_counts()
room_count_x_output_label = X_train.groupby(['rooms', 'output_label']).Id.count().unstack()
room_distribution.plot(ax=axes[0], kind='pie')
room_count_x_output_label.plot(ax=axes[1], kind='bar')
```

Out[48]:

&lt;AxesSubplot:xlabel='rooms'&gt;



The graph above shows that most of the house holds have 5 rooms or 4 rooms per house of family covering more than 50% of data available in training set, however couldn't find a reliable pattern with output label

In [49]:

```
x_train.groupby(['output_label', 'v18q']).Id.count().unstack()
```

Out[49]:

|              | v18q | 0    | 1 |
|--------------|------|------|---|
| output_label |      |      |   |
| 1            | 561  | 44   |   |
| 2            | 1162 | 94   |   |
| 3            | 851  | 113  |   |
| 4            | 3324 | 1496 |   |

Above matrix shows that it is more likely that if a person owns a tablet, he is more likely to be category 4 compared to other categories but this could still also be due to the bias in data from category 4

we can make gender column from male and female column

In [50]:

```
x_train[['male', 'female']].head(5)
```

Out[50]:

|      | male | female |
|------|------|--------|
| 9025 | 0    | 1      |
| 742  | 0    | 1      |
| 180  | 1    | 0      |
| 1115 | 0    | 1      |
| 9090 | 1    | 0      |

In [51]:

```
x_train['gender'] = np.where(x_train['female'], 1, 0)
```



In [52]:

```
X_train.groupby(['output_label', 'gender']).Id.count().unstack()
```

Out[52]:

| gender       |  | 0    | 1    |
|--------------|--|------|------|
| output_label |  |      |      |
| 1            |  | 264  | 341  |
| 2            |  | 600  | 656  |
| 3            |  | 454  | 510  |
| 4            |  | 2406 | 2414 |

The distribution of poverty marking seems to qually distributed across genders

Similarly we can reduce Zone column into single column

In [53]:


```
X_train['zone'] = np.where(X_train['areal'], 1, 0)
```

In [54]:

```
X_train.groupby(['output_label', 'zone']).Id.count().unstack().plot(kind='bar', figs
```

Out[54]:

<AxesSubplot:xlabel='output\_label'>



Above graph shows a skewness of higher chance of people being in urban zone compared to rural and fall under category 4 however even for rural category 4 seems to be the prominent one.

Combining regions to get better perspective of distribution across region for poverty

In [55]:

```
X_train[['lugar1', 'lugar2', 'lugar3', 'lugar4', 'lugar5', 'lugar6']].head(3)
```

Out[55]:

|      | lugar1 | lugar2 | lugar3 | lugar4 | lugar5 | lugar6 |
|------|--------|--------|--------|--------|--------|--------|
| 9025 | 0      | 0      | 0      | 0      | 0      | 1      |
| 742  | 1      | 0      | 0      | 0      | 0      | 0      |
| 180  | 1      | 0      | 0      | 0      | 0      | 0      |

In [56]:

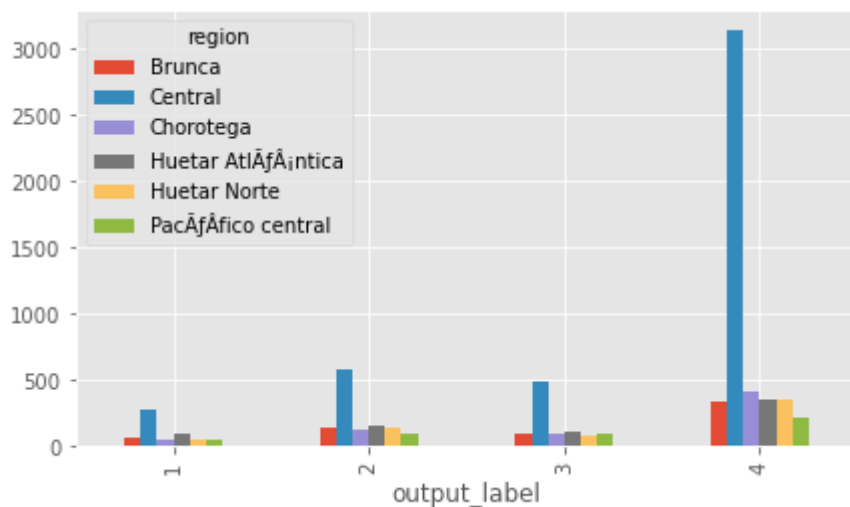
```
X_train['region'] = np.where(
    X_train['lugar1'] == 1, "Central",
    np.where(
        X_train['lugar2'] == 1, "Chorotega",
        np.where(
            X_train['lugar3'] == 1, "PacÃfÃfico central",
            np.where(
                X_train['lugar4'] == 1, "Brunca",
                np.where(
                    X_train['lugar5'] == 1, "Huetar AtlÃfÃintica",
                    "Huetar Norte"
                )
            )
        )
    )
)
```

In [57]:

```
X_train.groupby(['output_label', 'region']).Id.count().unstack().plot(kind='bar', fi
```

Out[57]:

&lt;AxesSubplot:xlabel='output\_label'&gt;



Above graph shows that majority of population is from central part and that too has a prominent category 4

population, this could also be due to inherit bias in data for category 4

In [58]:

```
region_map = {  
    "Central": 0, "Chorotega": 1, "PacÃfÃfico central": 2, "Brunca": 3, "Huetar AtlÃ
```

In [59]:


```
X_train['region'] = X_train['region'].map(region_map)
```

In [60]:

```
X_train['age'].plot(kind='hist', bins=100)
```

Out[60]:

<AxesSubplot:ylabel='Frequency'>




In [61]:

```
X_train['agesq'].plot(kind='hist')
```

Out[61]:

<AxesSubplot:ylabel='Frequency'>



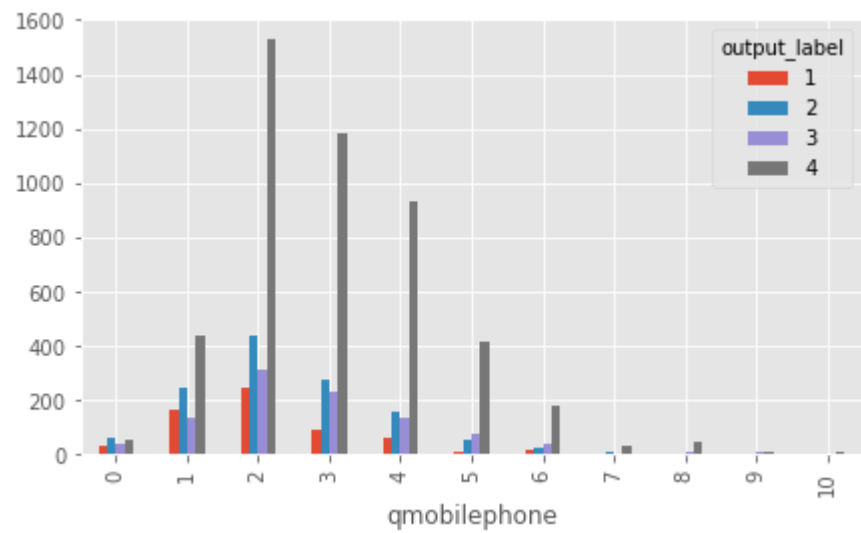
Age seems to be a right skewed distribution with majority of people being in age band of 15-30 years roughly.

In [62]:

```
X_train.groupby(['qmobilephone', 'output_label']).Id.count().unstack().plot(kind='bar')
```

Out[62]:

<AxesSubplot:xlabel='qmobilephone'>



In [63]:

```
X_train['has_mobile_or_phone'] = np.where(X_train['qmobilephone'] == 0, 0, 1)
```

In [64]:

```
X_train.groupby(['has_mobile_or_phone', 'output_label']).Id.count().unstack()
```

Out[64]:

| output_label        |   | 1   | 2    | 3   | 4    |
|---------------------|---|-----|------|-----|------|
| has_mobile_or_phone |   |     |      |     |      |
| 0                   | 0 | 31  | 60   | 39  | 54   |
|                     | 1 | 574 | 1196 | 925 | 4766 |

Couldn't find a valuable insight from a person being in poverty and having of mobile or phone

In [65]:

```
X_train[['tipovivi1', 'tipovivi2', 'tipovivi3', 'tipovivi4', 'tipovivi5']]
```

Out[65]:

|      | tipovivi1 | tipovivi2 | tipovivi3 | tipovivi4 | tipovivi5 |
|------|-----------|-----------|-----------|-----------|-----------|
| 9025 | 1         | 0         | 0         | 0         | 0         |
| 742  | 0         | 1         | 0         | 0         | 0         |
| 180  | 1         | 0         | 0         | 0         | 0         |
| 1115 | 0         | 0         | 1         | 0         | 0         |
| 9090 | 0         | 1         | 0         | 0         | 0         |
| ...  | ...       | ...       | ...       | ...       | ...       |
| 5734 | 0         | 1         | 0         | 0         | 0         |
| 5191 | 1         | 0         | 0         | 0         | 0         |
| 5390 | 0         | 0         | 0         | 0         | 1         |
| 860  | 0         | 0         | 1         | 0         | 0         |
| 7270 | 0         | 1         | 0         | 0         | 0         |

7645 rows × 5 columns

In [66]:

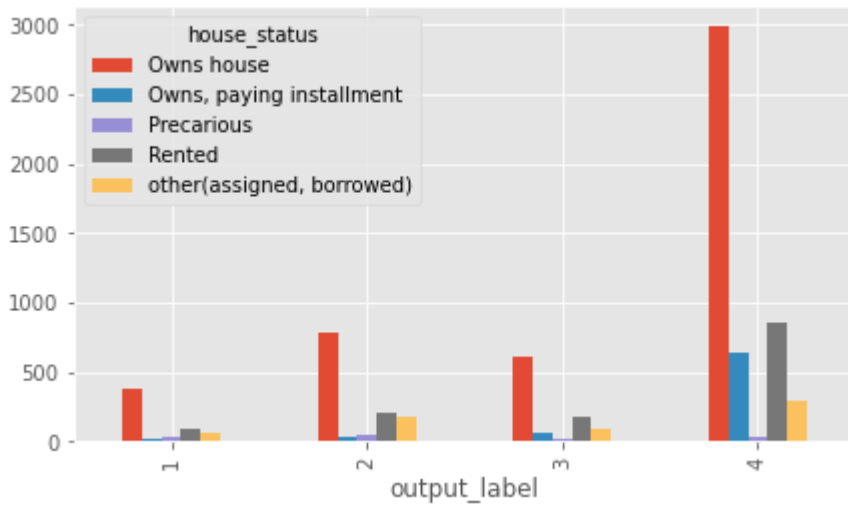
```
X_train['house_status'] = np.where(
    X_train['tipovivi1'] == 1, "Owns house",
    np.where(
        X_train['tipovivi2'] == 1, "Owns, paying installment",
        np.where(
            X_train['tipovivi3'] == 1, "Rented",
            np.where(
                X_train['tipovivi4'] == 1, "Precarious",
                "other(assigned, borrowed)"
            )
        )
    )
)
```

In [67]:

```
X_train.groupby(['output_label', 'house_status']).Id.count().unstack().plot(kind='bar')
```

Out[67]:

```
<AxesSubplot: xlabel='output_label'>
```



House status shows an interesting insight that you are less likely to be in category 4 of target value if your house status is Precarious

In [68]:

```
map_house_status = {
    "Owns house": 0, "Owns, paying installment": 1, "Rented": 2, "Precarious": 3, "other(assigned, borrowed)": 4
}
```

In [69]:

```
X_train['house_status'] = X_train['house_status'].map(map_house_status)
```

In [70]:

```
X_train = X_train.drop(columns = ['Id', 'idhogar', 'output_label'])
```

## Feature Engineering

In [71]:

```

# Gender
X_test['gender'] = np.where(X_test['female'], 1, 0)
# Zone
X_test['zone'] = np.where(X_test['areal'], 1, 0)
#Has mobile or not
X_test['has_mobile_or_phone'] = np.where(X_test['qmobilephone'] == 0, 0, 1)
# Region
X_test['region'] = np.where(
    X_test['lugar1'] == 1, "Central",
    np.where(
        X_test['lugar2'] == 1, "Chorotega",
        np.where(
            X_test['lugar3'] == 1, "PacÃfÃfico central",
            np.where(
                X_test['lugar4'] == 1, "Brunca",
                np.where(
                    X_test['lugar5'] == 1, "Hueta AtlÃfÃntica",
                    "Hueta Norte"
                )
            )
        )
    )
)

# House Status
X_test['house_status'] = np.where(
    X_test['tipovivil'] == 1, "Owns house",
    np.where(
        X_test['tipovivi2'] == 1, "Owns, paying installment",
        np.where(
            X_test['tipovivi3'] == 1, "Rented",
            np.where(
                X_test['tipovivi4'] == 1, "Precarious",
                "other(assigned, borrowed)"
            )
        )
    )
)

```

In [72]:

```
X_test['region'] = X_test['region'].map(region_map)
```

In [73]:

```
X_test['house_status'] = X_test['house_status'].map(map_house_status)
```

In [74]:

```
X_test = X_test.drop(columns = ['Id', 'idhogar'])
```

## Models

### RandomForest

In [75]:

```
model = RandomForestClassifier()
```

In [76]:

```
model.fit(X_train,y_train)
```

```
/var/folders/3c/lwghlsts08s19tyjdbv2ssd00000gn/T/ipykernel_93407/27213
49307.py:1: DataConversionWarning: A column-vector y was passed when a
1d array was expected. Please change the shape of y to (n_samples,), f
or example using ravel().
```

```
model.fit(X_train,y_train)
```

Out[76]:

```
RandomForestClassifier()
```

In [77]:

```
y_predict = model.predict(X_test)
```

```
/Users/saishruthicherukuri/opt/anaconda3/lib/python3.9/site-packages/s
klearn/base.py:493: FutureWarning: The feature names should match thos
e that were passed during fit. Starting version 1.2, an error will be
raised.
```

```
Feature names must be in the same order as they were in fit.
```

```
warnings.warn(message, FutureWarning)
```

In [78]:

```
print(accuracy_score(y_test,y_predict))
print(confusion_matrix(y_test,y_predict))
print(classification_report(y_test,y_predict))
```

```
0.9288702928870293
```

```
[[ 144    3    0   22]
 [   5  258    0   39]
 [   2    4  195   56]
 [   0    4    1 1179]]
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1            | 0.95      | 0.85   | 0.90     | 169     |
| 2            | 0.96      | 0.85   | 0.90     | 302     |
| 3            | 0.99      | 0.76   | 0.86     | 257     |
| 4            | 0.91      | 1.00   | 0.95     | 1184    |
| accuracy     |           |        | 0.93     | 1912    |
| macro avg    | 0.95      | 0.87   | 0.90     | 1912    |
| weighted avg | 0.93      | 0.93   | 0.93     | 1912    |

### Random Forest with K Folds

In [79]:

```
from sklearn.model_selection import KFold,cross_val_score
```



In [80]:

```
x_features=data_train.iloc[:,0:-1]
y_features=data_train.iloc[:,-1]
print(x_features.shape)
print(y_features.shape)
```

```
(9557, 140)
(9557,)
```

In [81]:

```
x_features = x_features.drop(columns = ['Id', 'idhogar'])
```

### **Checking the score using default 10 trees**

In [82]:

```
seed=7
kfold=KFold(n_splits=5,random_state=seed,shuffle=True)
model=RandomForestClassifier(random_state=10,n_jobs = -1)
print(cross_val_score(model,x_features,y_features,cv=kfold,scoring='accuracy',error_
results=cross_val_score(model,x_features,y_features,cv=kfold,scoring='accuracy',erro
print(results.mean()*100)
```

```
[0.93410042 0.94142259 0.92935636 0.93040293 0.93092622]
93.32417035047041
```

### **Checking the score using 100 trees**

In [83]:

```
num_trees= 100

model=RandomForestClassifier(n_estimators=100, random_state=10,n_jobs = -1)
print(cross_val_score(model,x_features,y_features,cv=kfold,scoring='accuracy'))
results=cross_val_score(model,x_features,y_features,cv=kfold,scoring='accuracy')
print(results.mean()*100)
```

```
[0.93410042 0.94142259 0.92935636 0.93040293 0.93092622]
93.32417035047041
```

In [84]:

```

model.fit(x_features,y_features)
labels = list(x_features)
feature_importances = pd.DataFrame({'feature': labels, 'importance': model.feature_i
feature_importances=feature_importances[feature_importances.importance>0.015]
feature_importances.head()

```

Out[84]:

|    | feature | importance |
|----|---------|------------|
| 0  | v2a1    | 0.015878   |
| 2  | rooms   | 0.020800   |
| 9  | r4h2    | 0.017319   |
| 10 | r4h3    | 0.016323   |
| 13 | r4m3    | 0.015872   |

In [85]:

```


feature_importances.sort_values(by=['importance'], ascending=True, inplace=True)
feature_importances['positive'] = feature_importances['importance'] > 0
feature_importances.set_index('feature',inplace=True)
feature_importances.head()

feature_importances.importance.plot(kind='barh', figsize=(11, 6),color = feature_imp
plt.xlabel('Importance')

```

Out[85]:

Text(0.5, 0, 'Importance')



In [86]:

feature\_importances

Out[86]:

|                        | importance | positive |
|------------------------|------------|----------|
| feature                |            |          |
| <b>r4t2</b>            | 0.015845   | True     |
| <b>r4m3</b>            | 0.015872   | True     |
| <b>v2a1</b>            | 0.015878   | True     |
| <b>r4h3</b>            | 0.016323   | True     |
| <b>r4t1</b>            | 0.016652   | True     |
| <b>cielorazo</b>       | 0.016670   | True     |
| <b>r4h2</b>            | 0.017319   | True     |
| <b>SQBhogar_nin</b>    | 0.019292   | True     |
| <b>hogar_nin</b>       | 0.019557   | True     |
| <b>rooms</b>           | 0.020800   | True     |
| <b>overcrowding</b>    | 0.023703   | True     |
| <b>SQBedjefe</b>       | 0.024987   | True     |
| <b>SQBovercrowding</b> | 0.025591   | True     |
| <b>qmobilephone</b>    | 0.025620   | True     |
| <b>SQBdependency</b>   | 0.027054   | True     |
| <b>dependency</b>      | 0.031348   | True     |
| <b>meaneduc</b>        | 0.051477   | True     |
| <b>SQBmeaned</b>       | 0.054520   | True     |

In [87]:

```
for i in feature_importances:
    if i not in x_features:
        print(i)
```

importance  
positive

In [88]:

```
X_data_Top_features= x_features[ ['r4t2','r4m3','v2a1','r4h3','r4t1','cielorazo','r4h2',
                                   'SQBhogar_nin','hogar_nin','rooms','overcrowding','SQBedjefe',
                                   'qmobilephone','SQBdependency','dependency','meaneduc']]
```

In [89]:

```
X_train,X_test,Y_train,Y_test=train_test_split(X_data_Top_features,
                                              y_features,
                                              test_size=0.2,
                                              stratify=y_features,random_state=0)
```

In [90]:

```
X_data_Top_features.isnull().sum()
```

Out[90]:

|                 |       |
|-----------------|-------|
| r4t2            | 0     |
| r4m3            | 0     |
| v2a1            | 0     |
| r4h3            | 0     |
| r4t1            | 0     |
| cielorazo       | 0     |
| r4h2            | 0     |
| r4t1            | 0     |
| cielorazo       | 0     |
| r4h2            | 0     |
| SQBhogar_nin    | 0     |
| hogar_nin       | 0     |
| rooms           | 0     |
| overcrowding    | 0     |
| SQBedjefe       | 0     |
| SQBovercrowding | 0     |
| qmobilephone    | 0     |
| SQBdependency   | 0     |
| dependency      | 0     |
| meaneduc        | 0     |
| SQBmeaned       | 0     |
| dtype:          | int64 |

In [91]:

```
model = RandomForestClassifier()
model_final=model.fit(X_train,Y_train)
Y_pred=model_final.predict(X_test)
```

In [92]:

```
print(accuracy_score(Y_test,Y_pred))
print(confusion_matrix(Y_test,Y_pred))
print(classification_report(Y_test,Y_pred))
print(f1_score(Y_test,Y_pred,average='weighted'))
```

```
0.9633891213389121
[[ 144    0    0   11]
 [   6  292    3   11]
 [   1    3  228   12]
 [   2    7   14 1178]]
      precision    recall  f1-score   support

     1         0.94      0.93      0.94         155
     2         0.97      0.94      0.95         312
     3         0.93      0.93      0.93         244
     4         0.97      0.98      0.98        1201

 accuracy                   0.96         1912
 macro avg              0.95      0.95      0.95         1912
 weighted avg           0.96      0.96      0.96         1912
```

0.9633129858012178

## Test Data Cleaning and Prediction

In [93]:

```
# lets drop required variables.
data_test.drop(['Id', 'idhogar', 'edjefe', 'edjefa'],axis=1,inplace=True)
data_test['dependency']=data_test['dependency'].apply(map)
```

In [94]:

```
data_test['v2a1'].fillna(0,inplace=True)
data_test['v18q1'].fillna(0,inplace=True)
data_test['SQBmeanead'].fillna(0,inplace=True)
data_test['meaneaduc'].fillna(0,inplace=True)
```

In [95]:

```
data_test.isna().sum().value_counts()
```

Out[95]:

```
0    138
dtype: int64
```

In [96]:

```
data_test=data_test[['r4t2', 'r4m3', 'v2a1', 'r4h3', 'r4t1', 'cielorazo', 'r4h2', 'r4t1', 'c
SQBhogar_nin', 'hogar_nin', 'rooms', 'overcrowding', 'SQ
'qmobilephone', 'SQBdependency', 'dependency', 'meaneadu
```

In [97]:

```
test_prediction=model_final.predict(data_test)
```

In [98]:

```
test_prediction
```

Out[98]:

```
array([4, 4, 4, ..., 2, 2, 2])
```

***Using RandomForest Classifier we can predict test data with accuracy of 96%.***

In [ ]: